

[◀ Return to Classroom](#)

Deploying a Sentiment Analysis Model

REVIEW

CODE REVIEW 3

HISTORY

▶ website/index.html 1

▶ train/train.py 1

▼ serve/predict.py 1

```
1 import argparse
2 import json
3 import os
4 import pickle
5 import sys
6 import sagemaker_containers
7 import pandas as pd
8 import numpy as np
9 import torch
10 import torch.nn as nn
11 import torch.optim as optim
12 import torch.utils.data
13
14 from model import LSTMClassifier
15
16 from utils import review_to_words, convert_and_pad
17
18 def model_fn(model_dir):
19     """Load the PyTorch model from the `model_dir` directory."""
20     print("Loading model.")
```

```

21 # First, load the parameters used to create the model.
22 model_info = {}
23 model_info_path = os.path.join(model_dir, 'model_info.pth')
24 with open(model_info_path, 'rb') as f:
25     model_info = torch.load(f)
26
27
28 print("model_info: {}".format(model_info))
29
30 # Determine the device and construct the model.
31 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
32 model = LSTMClassifier(model_info['embedding_dim'], model_info['hidden_dim'], model_info['output_dim'])
33
34 # Load the store model parameters.
35 model_path = os.path.join(model_dir, 'model.pth')
36 with open(model_path, 'rb') as f:
37     model.load_state_dict(torch.load(f))
38
39 # Load the saved word_dict.
40 word_dict_path = os.path.join(model_dir, 'word_dict.pkl')
41 with open(word_dict_path, 'rb') as f:
42     model.word_dict = pickle.load(f)
43
44 model.to(device).eval()
45
46 print("Done loading model.")
47 return model
48
49 def input_fn(serialized_input_data, content_type):
50     print('Deserializing the input data.')
51     if content_type == 'text/plain':
52         data = serialized_input_data.decode('utf-8')
53         return data
54     raise Exception('Requested unsupported ContentType in content_type: ' + content_type)
55
56 def output_fn(prediction_output, accept):
57     print('Serializing the generated output.')
58     return str(prediction_output)
59
60 def predict_fn(input_data, model):
61     print('Inferring sentiment of input data.')
62
63     device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
64
65     if model.word_dict is None:
66         raise Exception('Model has not been loaded properly, no word_dict.')
67
68     # TODO: Process input_data so that it is ready to be sent to our model.
69     # You should produce two variables:
70     # data_X - A sequence of length 500 which represents the converted review
71     # data_len - The length of the review
72
73     #data_X = None
74     #data_len = None
75
76     words = review_to_words(input_data)
77     data_X, data_len = convert_and_pad(model.word_dict, words)
78
79     # Using data_X and data_len we construct an appropriate input tensor. Remember
80     # that our model expects input data of the form 'len, review[500]'.
81     data_pack = np.hstack((data_len, data_X))
82     data_pack = data_pack.reshape(1, -1)

```

```
83
84 data = torch.from_numpy(data_pack)
85 data = data.to(device)
86
87 # Make sure to put the model into evaluation mode
88 model.eval()
89
90 # TODO: Compute the result of applying the model to the input data. The variable
91 #       be a numpy array which contains a single integer which is either 1 or 0
92
93 with torch.no_grad():
94     out = model.forward(data)
95
96 result = np rint(out.numpy())
```

AWESOME

Expressions correct

```
97
98
99 return result
100
```

RETURN TO PATH